

Further Reading

 coursera.org/learn/uol-machine-learning-for-all/supplement/vSt7N/further-reading

There has now been a lot written about Machine Learning and AI and their societal effects. Here are some good places to start:

Domingos, Pedro (2015). *The Master Algorithm: How the Quest for the Ultimate Learning Machine Will Remake Our World*. New York: Basic Books. ISBN 978-0-465-06570-7.

This is my favourite beginners book on machine learning. It is a really good explanation of the techniques and approaches for someone without a computer science background.

Cathy O'Neil. (2016) *Weapons of Math Destruction: How Big Data Increases Inequality and Threatens Democracy*. Crown Publishing Group, New York, NY, USA. ISBN:0553418815 9780553418811

This is an excellent book on the dangers of machine learning in modern life and how data can have unexpected negative effects.

Kate Crawford and Ryan Calo, (2016) 'There is a blind spot in AI research', Nature, October 13. <http://www.nature.com/news/there-is-a-blind-spot-in-ai-research-1.20805>

Kate Crawford is one of the most important thinkers on the dangers of AI, she hasn't written a book, but this article is excellent.

Bostrom, Nick (2014). *Superintelligence: Paths, Dangers, Strategies*. Oxford: Oxford University Press. pp. 229–237. ISBN 978-0199678112.

Probably the most famous book is Superintelligence, by Nick Bostrom, though I would say that this focuses too much on the theoretical future impacts of AI, while there are a lot of challenges that we need to deal with right now.